


# R language and data analysis: data manipulation advanced


Qiang Shen

Dec. 17, 2017

# data analysis procedure



# data analysis procedure



# A quick demo 1

- split-apply-combine in base

```
with(iris, tapply(iris[, 1], Species,  
  mean))
```

setosa	versicolor	virginica
5.006	5.936	6.588

## A quick demo 2

- split-apply-combine in base

```
s <- split(iris, iris$Species)
tmp.data <- iris[, 1:3]
# head(tmp.data)
sapply(s, function(x) colMeans(tmp.data,
  na.rm = T))
```

	setosa	versicolor
Sepal.Length	5.843333	5.843333
Sepal.Width	3.057333	3.057333
Petal.Length	3.758000	3.758000

	virginica
Sepal.Length	5.843333
Sepal.Width	3.057333
Petal.Length	3.758000

# dplyr

```
library(dplyr)
iris %>% group_by(Species) %>% select(one_of("Sepal.Length",
      summarise_each(funs(mean(.))))
iris %>% group_by(Species) %>% select(one_of(names(iris[,
      1:3]))) %>% summarise_each(funs(mean(.)))
```

```
# A tibble: 3 x 2
  Species Sepal.Length
  <fctr>      <dbl>
1   setosa      5.006
2 versicolor  5.936
3  virginica   6.588
```

```
# A tibble: 3 x 4
  Species Sepal.Length
  <fctr>      <dbl>
1   setosa      5.006
2 versicolor  5.936
3  virginica   6.588
# ... with 2 more variables:
#   Sepal.Width <dbl>,
#   Petal.Length <dbl>
```

# dplyr: tibble

```
library(tibble)
iris
head(iris)
as.tibble(iris)
tbl_df(iris)
```



# dplyr: tibble

```
tibble(x = 1:5, y = 1, z = x^2 + y)
```

```
# A tibble: 5 x 3
```

	x	y	z
	<int>	<dbl>	<dbl>
1	1	1	2
2	2	1	5
3	3	1	10
4	4	1	17
5	5	1	26

## dplyr: tibble

```
tribble(~x, ~y, ~z, "a", 2, 3.6, "b",  
        1, 8.5)
```

```
# A tibble: 2 x 3  
      x     y     z  
  <chr> <dbl> <dbl>  
1     a     2   3.6  
2     b     1   8.5
```

## dplyr: tibble

```
df1 <- data.frame(x = 1:3, y = 3:1)  
class(df1[, 1:2])
```

```
[1] "data.frame"
```

```
class(df1[, 1])
```

```
[1] "integer"
```

```
df2 <- tibble(x = 1:3, y = 3:1)  
class(df2[, 1:2])
```

```
[1] "tbl_df"      "tbl"  
[3] "data.frame"
```

```
class(df2[, 1])
```

# dplyr features

- a powerful R-package to transform and summarize tabular data with rows and columns.
- ① intuitive to translate your thoughts into codes/scripts.
  - ② spend less time waiting for the computer.

# key verbs in dplyr package

- `filter()` (and `slice()`)
- `arrange()`
- `select()` (and `rename()`)
- `distinct()`
- `mutate()` (and `transmute()`)
- `summarise()`
- `sample_n()` and `sample_frac()`

## key verbs

- filter: keep rows matching criteria
- select: pick columns by name
- arrange: reorder rows
- mutate: add new variables
- summarise: reduce variables to values

# key verbs

## Data frame



Rows: filter



Columns: select



mutate(\_each)

transmute



summarise(\_each)

## structure of verbs

- First argument is a data frame, say *flights*.
- Subsequent arguments say what to do with the data frame.
- The result is a data frame



# basics of dataset

- dataset contains all 336776 flights that departed from New York City in 2013.

```
library(dplyr)
library(nycflights13)
dim(flights)
head(flights)
flights  #tbl_df
print(tbl_df(mtcars), n = 5)
tbl_df(iris)
glimpse(flights)
```

# key verbs

- **filter()** (and **slice()**)
- **arrange()**
- **select()** (and **rename()**)
- **distinct()**
- **mutate()** (and **transmute()**)
- **summarise()**
- **sample\_n()** and **sample\_frac()**

# base function

- all flights on January 1st



```
flights[flights$month == 1 & flights$day ==  
1, ]
```

## base function

- all flights on January 1st

```
library(nycflights13)  
subset(flights, month == 1 & day == 1)
```

# filter()

```
library(nycflights13)
filter(flights, month == 1, day == 1)  ##tidyr ':' vs. ',
filter(flights, month == 1 | month ==
      2)
filter(flights, carrier == "AA" | carrier ==
      "UA")
filter(flights, carrier %in% c("AA",
      "UA"))
```

# logic for filter()

Logic in R - ?Comparison, ?base::Logic			
<	Less than	!=	Not equal to
>	Greater than	%in%	Group membership
==	Equal to	is.na	Is NA
<=	Less than or equal to	!is.na	Is not NA
>=	Greater than or equal to	&,  , !, xor, any, all	Boolean operators

## select rows by position: base

```
library(dplyr)
library(nycflights13)
flights[1:10, ]
```

## select rows by position: slice

```
library(nycflights13)  
slice(flights, 1:10)
```



## key verbs

- `filter()` (and `slice()`)
- **`arrange()`**
- `select()` (and `rename()`)
- `distinct()`
- `mutate()` (and `transmute()`)
- `summarise()`
- `sample_n()` and `sample_frac()`

## order data: base

- Order data based on specified columns.

```
flights[order(flights$year, flights$month,  
             flights$day), ]  
flights[order(-flights$arr_delay), ]  #descend
```

## order data: arrange

- Order data based on specified columns.

```
library(nycflights13)
arrange(flights, year, month, day)
arrange(flights, desc(arr_delay))
```

## key verbs

- `filter()` (and `slice()`)
- `arrange()`
- **`select()` (and `rename()`)**
- `distinct()`
- `mutate()` (and `transmute()`)
- `summarise()`
- `sample_n()` and `sample_frac()`

# Select columns by name: base

- select columns of year, month and age



```
library(nycflights13)
flights[, c("year", "month", "day")]
subset(flights, select = (year:day))
```

## Select columns by name: select

- Select columns between year and day (inclusive)

```
library(nycflights13)
select(flights, year, month, day)
select(flights, year:day)
```

## Select columns by name: select

- Select all columns except those from year to day (inclusive)

```
library(nycflights13)  
select(flights, -(year:day))
```

## select:rename

```
library(nycflights13)  
select(flights, tail_num = tailnum)
```



## verb rename in dplyr

- rename a variable

```
library(nycflights13)
rename(flights, tail_num = tailnum)
flights
```

## key verbs

- `filter()` (and `slice()`)
- `arrange()`
- `select()` (and `rename()`)
- **`distinct()`**
- `mutate()` (and `transmute()`)
- `summarise()`
- `sample_n()` and `sample_frac()`

# distinct vs. unique

IDs
1
1
1
2
2
2
4
4
4
8
8
8
12
12
12

base::unique



dplyr::distinct

IDs
1
2
4
8
12

## distinct vs. unique

```
library(nycflights13)
unique(select(flights, tailnum))
distinct(select(flights, tailnum))
distinct(select(flights, origin, dest))
```

## key verbs: new column

- `filter()` (and `slice()`)
- `arrange()`
- `select()` (and `rename()`)
- `distinct()`
- **`mutate()` (and `transmute()`)**
- `summarise()`
- `sample_n()` and `sample_frac()`

## new column: base



```
iris$Sepal_sum <- iris$Sepal.Length +  
  iris$Sepal.Width  
head(iris)[, c(1:2, 5:6)]
```

	Sepal.Length	Sepal.Width	Species
1	5.1	3.5	setosa
2	4.9	3.0	setosa
3	4.7	3.2	setosa
4	4.6	3.1	setosa
5	5.0	3.6	setosa
6	5.4	3.9	setosa

## new column: transform in base

```
library(nycflights13)
transform(flights, gain = arr_delay -
  dep_delay, speed = distance/air_time *
  60)
```

## create new coulmun

### Make New Variables



**dplyr::mutate(iris, sepal = Sepal.Length + Sepal. Width)**

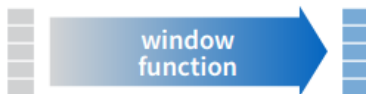
Compute and append one or more new columns.

**dplyr::mutate\_each(iris, funs(min\_rank))**

Apply window function to each column.

**dplyr::transmute(iris, sepal = Sepal.Length + Sepal. Width)**

Compute one or more new columns. Drop original columns.





## new column: mutate

```
library(nycflights13)
mutate(flights, gain = arr_delay - dep_delay,
       speed = distance/air_time * 60)
```

## transform vs. mutate

```
mutate(flights, gain = arr_delay - dep_delay,  
       gain_per_hour = gain/(air_time/60))
```

```
transform(flights, gain = arr_delay -  
           delay, gain_per_hour = gain/(air_time/60))
```

# transmute

```
transmute(flights, gain = arr_delay -  
  dep_delay, gain_per_hour = gain/(air_time/60))
```

## mutate\_each

```
iris_new <- mutate_each(iris[, 1:4],  
  funs(ratio = ./max(.), avg = mean(.)))  
head(iris_new)
```

# key verbs in dplyr package

- `filter()` (and `slice()`)
- `arrange()`
- `select()` (and `rename()`)
- `distinct()`
- `mutate()` (and `transmute()`)
- **`summarise()`**
- `sample_n()` and `sample_frac()`

## summarise



```
library(nycflights13)
summarise(flights, delay = mean(dep_delay,
  na.rm = TRUE))
```

```
# A tibble: 1 x 1
  delay
  <dbl>
1 12.63907
```

## Summarise Data



**dplyr::summarise(iris, avg = mean(Sepal.Length))**

Summarise data into single row of values.

**dplyr::summarise\_each(iris, funs(mean))**

Apply summary function to each column.

**dplyr::count(iris, Species, wt = Sepal.Length)**

Count number of rows with each unique value of variable (with or without weights).



# summarise



Summarise uses **summary functions**, functions that take a vector of values and return a single value, such as:

## **dplyr::first**

First value of a vector.

## **dplyr::last**

Last value of a vector.

## **dplyr::nth**

Nth value of a vector.

## **dplyr::n**

# of values in a vector.

## **dplyr::n\_distinct**

# of distinct values in a vector.

## **IQR**

IQR of a vector.

## **min**

Minimum value in a vector.

## **max**

Maximum value in a vector.

## **mean**

Mean value of a vector.

## **median**

Median value of a vector.

## **var**

Variance of a vector.

## **sd**

Standard deviation of a vector.



## summarise\_each

```
library(nycflights13)
summarise_each(iris[, 1:4], funs(avg = mean(.,
  na.rm = T)))
```

	Sepal.Length_avg	Sepal.Width_avg
1	5.843333	3.057333
	Petal.Length_avg	Petal.Width_avg
1	3.758	1.199333

## count

```
# count(iris, Species)
count(iris, Species, wt = Sepal.Length)
```

```
# A tibble: 3 x 2
  Species      n
  <fctr> <dbl>
1  setosa 250.3
2 versicolor 296.8
3  virginica 329.4
```

```
with(iris, tapply(iris[, 1], Species,
  sum))
```

setosa	versicolor	virginica
250.3	296.8	329.4

# key verbs in dplyr package

- `filter()` (and `slice()`)
- `arrange()`
- `select()` (and `rename()`)
- `distinct()`
- `mutate()` (and `transmute()`)
- `summarise()`
- **`sample_n()` and `sample_frac()`**

## Randomly sample rows.



```
sample_n(flights, 10)
sample_frac(flights, 0.01)
sample_frac(flights, 0.01, replace = T) ##bootstrap
```

# key verbs in dplyr package

- `filter()` (and `slice()`)
- `arrange()`
- `select()` (and `rename()`)
- `distinct()`
- `mutate()` (and `transmute()`)
- `summarise()`
- `sample_n()` and `sample_frac()`

## key verbs

- filter: keep rows matching criteria
- select: pick columns by name
- mutate: add new variables
- summarise: reduce variables to values

# key verbs

## Data frame



Rows: filter



Columns: select



mutate

transmute



summarise

## structure of verbs

- The first argument is a data frame. (import)
- The subsequent arguments describe what to do with it.
- The result is a new data frame. (export)



Practice!

# group\_by

## Group Data

**dplyr::group\_by(iris, Species)**


Group data into rows with the same value of Species.

**dplyr::ungroup(iris)**

Remove grouping information from data frame.


**iris %>% group\_by(Species) %>% summarise(...)**

Compute separate summary row for each group.



**iris %>% group\_by(Species) %>% mutate(...)**

Compute new variables by group.



## group\_by

```
library(nycflights13)
by_tailnum <- group_by(flights, tailnum)
by_tailnum
delay <- summarise(by_tailnum, count = n(),
  dist = mean(distance, na.rm = TRUE),
  delay = mean(arr_delay, na.rm = TRUE))
delay <- filter(delay, count > 20, dist <
  2000)
with(delay, plot(dist, delay))
abline(with(delay, lm(delay ~ dist)))
```

# aggregate functions

- package base: `min()`, `max()`, `mean()`, `sum()`, `sd()`, `median()`, and `IQR()`.
- package dplyr: `n()`, `n_distinct(x)`, `first(x)`, `last(x)` and `nth(x, n)`

```
destinations <- group_by(flights, dest)
destinations
dest <- summarise(destinations, planes = n_distinct(tailnum),
  flights = n())
arrange(dest, desc(planes))
summarise(dest, sum(flights))
```

## aggregate functions

```
daily <- group_by(flights, month, day)
(per_day <- summarise(daily, flights = n()))
(per_month <- summarise(per_day, flights = sum(flights)))
(per_year <- summarise(per_month, flights = sum(flights)))
```

## group\_by: mutate with summarise

```
flights <- mutate(flights, hrs = air_time/60,  
  speed = distance/hrs)  
flights0 <- group_by(flights, tailnum)  
flight <- select(flights0, speed)  
flight  
summarise(flight, speed_mean = mean(speed,  
  na.rm = T))
```

## Chaining: example

```
a1 <- group_by(flights, year, month,
               day)
a2 <- select(a1, arr_delay, dep_delay)
a3 <- summarise(a2, arr = mean(arr_delay,
                               na.rm = TRUE), dep = mean(dep_delay,
                                                             na.rm = TRUE))
a4 <- filter(a3, arr > 30 | dep > 30)
```

## Chaining: example

```
filter(summarise(select(group_by(flights,  
  year, month, day), arr_delay, dep_delay),  
  arr = mean(arr_delay, na.rm = TRUE),  
  dep = mean(dep_delay, na.rm = TRUE)),  
  arr > 30 | dep > 30)
```



## Chaining: example

- %>%

```
flights %>% group_by(year, month, day) %>%  
  select(arr_delay, dep_delay) %>%  
  summarise(arr = mean(arr_delay, na.rm = TRUE),  
            dep = mean(dep_delay, na.rm = TRUE)) %>%  
  filter(arr > 30 | dep > 30)
```

## summary

- filter: keep rows matching criteria
- select: pick columns by name
- mutate: add new variables
- summarise: reduce variables to values
- chaining: %>%

